

Cherry Instruction Set (cisa)

16 bit instructions. 5 bits to decide which instruction to use. Remaining bits for parameters.

4 registers: input, weights, acc, out. Each holds a small matrix tile (4x4, 16x16, 32x32 depending on cherry device)

Arithmetic Instructions

| Instruction | Example | Meaning | Comments |
|----------------------------|--------------------|---|--|
| Matrix Multiply | matmul | REG[ACC] += REG[INPUT] @ REG[WEIGHTS] | |
| Multiply Accumulate | mulacc | REG[ACC] += REG[INPUT] * REG[WEIGHTS] | Element wise multiply |
| Vector Add | add | REG[OUTPUT] = REG[INPUT] + REG[WEIGHTS] | |
| Vector Subtract | sub | REG[OUTPUT] = REG[INPUT] - REG[WEIGHTS] | |
| Vector Multiply | mul | REG[OUTPUT] = REG[INPUT] * REG[WEIGHTS] | |
| Vector Divide | div | REG[OUTPUT] = REG[INPUT] / REG[WEIGHTS] | |
| Vector Power | pow | REG[OUTPUT] = REG[INPUT] ** REG[WEIGHTS] | |
| Vector Reduce Max | max \$1 \$2 \$3 | max([REG[ACC] if \$1 else none, (REG[OUT])]). \$2 is axis, \$3 is keep dims | Gets the max value in out reg or out reg and acc reg |
| Vector Reduce Sum | sum \$1, \$2, \$3 | sum([REG[ACC] if \$1 else none, (REG[OUT])], \$2 is axis, \$3 is count | Gets the sum of out reg or out reg + acc |
| Vector ReLU | relu | max(x,0) of every x in REG[INPUT] | |
| Vector Exp | exp | e^x of every x in REG[INPUT] | |
| Vector Log | log | log_2 of every x in REG[INPUT] | |
| Vector GT0 | gtz | 1 if x > 0 else 0 of every x in REG[INPUT] | |
| Copy ACC | copy_acc_to_output | REG[OUTPUT] = REG[ACC] | |
| Zero ACC | zero_acc | REG[ACC] = 0 | |

Data Transfer Instructions

| Instruction | Example | Meaning | Comments |
|------------------|-----------------------------------|---|---|
| Load | load \$1, \$2, \$3, \$4, \$5, \$6 | REG[\$2] = ... \$3 and \$4 are lengths. \$5 is zero flag, \$6 is skip flag. APU \$1 | Uses address from an APU. Cache to registers. |
| Store | store \$1, \$2, \$3, \$4, \$5 | MEM[addr] = REG[\$2][:\$3]:\$4 and APU \$1. If \$5 then skip local cache and DMA. Useful for when you need a value for backward pass but wont read it again from cache. | Use address from an APU. Registers to cache. |
| Mem Read | mem_read \$1, \$2, \$3 | \$1 and \$2 are cache addr \$3 is APU | ram to cache |
| Mem Write | mem_write \$1, \$2, \$3 | \$1 and \$2 are cache addr \$3 is APU | cache to ram |

Control Flow Instructions

| Instruction | Example | Meaning | Comments |
|-------------------------------|----------------------------------|-------------------------------------|---|
| start independent loop | start_independent_loop (\$1) | start loop with ro_data[\$1] params | From ro_data registers gets loop iteration count and number of instructions in the loop body. Loop iterations can be reordered freely. must be no loops nested inside this loop |
| start loop | start_loop (\$1) | start loop with ro_data[\$1] params | This loop must be executed in order. We can't rearrange when different loop bodies run. This is low performance loop |
| end loop | end_loop_or_jump_if_needed (\$1) | Gets jump number from ro_data[\$1] | jumps if loop has more more iterations remaining |

Read only data (program header)

Not all instructions can access all registers. Furthermore, read-only data is considered ro_data compiled into the kernel. This header section is 1212 bits.

- 8 loops allowed
 - Each loop needs 15 bits for the number of iterations
- 4 cisa_load/cisa_store address formulas allowed
 - Each address is a linear equation with 8 coefficients (gradients w.r.t. each loop variable) and an offset. Each value is 15 bits so one linear equation needs 135 bits
 - 30 bits for hard coded stride_x and stride_y
- 4 cisa_mem_* address formulas allowed
 - Same as cisa_load/cisa_store but now our addresses are bigger so the offset needs to be 18 bits instead of 15. So each linear equation needs 138 bits.

Thoughts

A program has 1332 bit ro_data in its header. On average 160 bits body for 10 instructions. So each program is 186 bytes.

128KB icache is 686 programs. A program is only valid for certain input tensor shapes and input tensor memory locations (cache or device memory)

Variable sized program headers can decrease pcache transistor usage.